

Libraries

```
library(tidyverse)
library(here)
library(brms)
library(broom)
library(tidybayes)
library(tolerance)
library(wesanderson)
library(patchwork)
library(gt)
```

As with any frequentist simulation looking at coverage of a p-value, we can start with the population parameters known. From there we simulate many draws from that population and calculate the desired statistic (in this case the tolerance interval limits). By comparing each calculated tolerance interval to the true population coverage, we can determine the fraction of simulations to do not meet the specified coverage. If that number, as a fraction of the total number of sims, is less than or equal to the p-value, then the tolerance interval procedure is working as intended and provided the specified coverage, on average, over many simulated experiments.

This block sets up the parameters and identifies the specified quantiles of the true population.

Simulated Experiments

True Population Parameters

```
set.seed(9989)

n_sims <- 10000
true_population_mean <- 40
true_population_sd <- 4
alpha = 0.1
p = .95

upper_97.5 <- qnorm(p = .975, mean = true_population_mean, sd = true_population_sd)
lower_02.5 <- qnorm(p = .025, mean = true_population_mean, sd = true_population_sd)
```

Now run the simulation, replicating many "experiments" of n=15, calculating a 90% confidence, 2-sided tolerance interval for 95% of the population each time.

Run the Simulation

```
sim_summary_tbl <- tibble(sim_id = seq(from = 1, to = n_sims, by = 1)) %>%
  rowwise() %>%
  mutate(
    test_sample_size = 15,
    sim_test_data = list(rnorm(n = test_sample_size, mean = true_population_mean, sd =
true_population_sd))
  ) %>%
  summarize(
    sim_id = sim_id,
    # sim_test_data = sim_test_data,
    norm_tol = list(normtol.int(sim_test_data, alpha = alpha, P = p, side = 2)),
    sample_sd = sd(sim_test_data)
  ) %>%
  unnest(norm_tol) %>%
  rename(
    lower_bound_ltl = "2-sided.lower",
    upper_bound_utl = "2-sided.upper",
    sample_mean = x.bar,
    p = P
  ) %>%
  ungroup() %>%
  select(sim_id, alpha, p, sample_mean, sample_sd, everything()) %>%
  mutate(
    true_lower_2.5 = lower_02.5,
    true_upper_97.5 = upper_97.5
  )

sim_summary_tbl %>%
  gt_preview()
```

	sim_id	alpha	p	sample_mean	sample_sd	lower_bound_ltl	upper_bound_utl	true_lower_2.5	true_upper_97.5
1	1	0.1	0.95	36.19320	3.243531	27.35493	45.03147	32.16014	47.83986
2	2	0.1	0.95	40.56155	3.766514	30.29821	50.82489	32.16014	47.83986
3	3	0.1	0.95	40.23154	4.268292	28.60091	51.86217	32.16014	47.83986
4	4	0.1	0.95	40.24057	4.175918	28.86165	51.61949	32.16014	47.83986
5	5	0.1	0.95	40.41020	3.880094	29.83737	50.98303	32.16014	47.83986
6..9999	10000	0.1	0.95	38.69646	4.780010	25.67145	51.72147	32.16014	47.83986

Visualize

Plot the resulting upper and lower bounds of the 95/90 tolerance interval. The red lines are the true population quantiles.

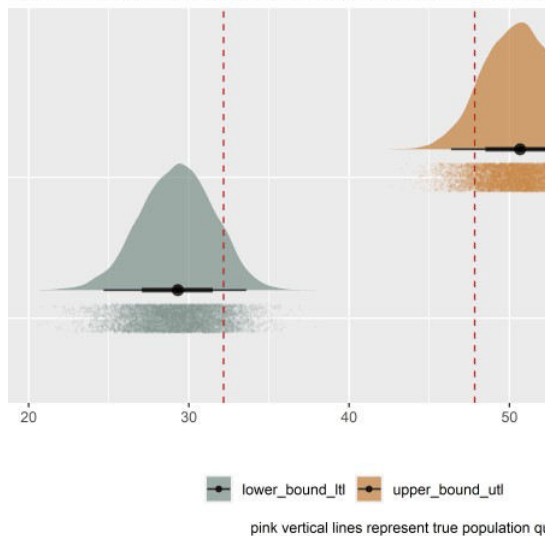
```
sim_summary_tbl %>%
  pivot_longer(cols = everything(), names_to = "param", values_to = "value") %>%
  filter(param == "lower_bound_ltl" | param == "upper_bound_utl") %>%
  mutate(param = as_factor(param)) %>%
  ggplot(aes(x = value, y = param)) +
```

```

geom_jitter(aes(color = param), width = .1, height = .1, size = .3, alpha = .05) +
stat_halfeye(aes(fill = param), alpha = .7, position = position_nudge(y = .2)) +
geom_vline(xintercept = mean(sim_summary_tbl$true_lower_2.5), color = "firebrick", linetype =
2) +
geom_vline(xintercept = mean(sim_summary_tbl$true_upper_97.5), color = "firebrick", linetype
= 2) +
scale_color_manual(values = wes_palette("Moonrise2")) +
scale_fill_manual(values = wes_palette("Moonrise2")) +
# scale_color_manual(values = c("purple", "limegreen")) +
# scale_fill_manual(values = c("purple", "limegreen")) +
labs(
  x = "",
  y = "",
  title = "Distribution of Calculated Upper and Lower Tolerance Limits",
  subtitle = "10,000 Simulations of n=15 parts with 2-sided normal Tolerance Intervals",
  caption = "pink vertical lines represent true population quantiles, .025 and .975"
) +
theme(legend.title = element_blank(),
      legend.position = "bottom")

```

Distribution of Calculated Upper and Lower Tolerance L
10,000 Simulations of n=15 parts with 2-sided normal Tolerance In



To see if the tolerance interval procedure is providing the desired coverage, we have to look at each set of lower and upper limits from each simulation. The code below converts the calculated tolerance bounds from each sim into quantiles of the true population, then calculates the difference to determine the true population coverage of the simulated experiment. If the coverage is less than intended, a flag of 1 is assigned in the "unacceptable_coverage" col, else 0.

Evaluate True Coverage of Each Simulated Experiment

```

coverage_summary_tbl <- sim_summary_tbl %>%
  mutate(ltl_coverage_true = map_dbl(lower_bound_ltl, pnorm, true_population_mean,
true_population_sd),
         utl_coverage_true = map_dbl(upper_bound_utl, pnorm, true_population_mean,
true_population_sd)) %>%
  mutate(total_tl_coverage_true = utl_coverage_true - ltl_coverage_true) %>%
  mutate(unacceptable_coverage = case_when(total_tl_coverage_true >= p ~ 0,
                                           TRUE ~ 1)) %>%
  select(-c(alpha, p, true_lower_2.5, true_upper_97.5))

```

```

coverage_summary_tbl %>%
  gt_preview()

```

	sim_id	sample_mean	sample_sd	lower_bound_ltl	upper_bound_utl	ltl_coverage_true	utl_coverage_true	total_tl_coverage_true	unacceptable_coverage
1	1	36.19320	3.243531	27.35493	45.03147	0.0007854236	0.8957802	0.8949948	1
2	2	40.56155	3.766514	30.29821	50.82489	0.0076447723	0.9965973	0.9889526	0
3	3	40.23154	4.268292	28.60091	51.86217	0.0021875221	0.9984893	0.9963017	0
4	4	40.24057	4.175918	28.86165	51.61949	0.0026797862	0.9981630	0.9954832	0
5	5	40.41020	3.880094	29.83737	50.98303	0.0055321969	0.9969814	0.9914492	0
6..9999									
10000	10000	38.69646	4.780010	25.67145	51.72147	0.0001703976	0.9983072	0.9981368	0

Count False Positives

For the coverage to be correct, the percentage of false positives (simulated experiments that covered less than p% of the true population) should be less than or equal to alpha (10%). That would mean less than $.1(10000) = 1000$ cases where the tolerance interval covered less than 95% of the true population.

```

coverage_summary_tbl %>%
  summarize(
    false_positives = sum(unacceptable_coverage),
    false_positive_limit = alpha * n_sims
  ) %>%
  mutate(
    alpha_level = alpha,

```

```

total_trials = n_sims,
false_positive_rate = false_positives / n_sims,
"tol_limit_procedure_working?" = case_when(
  false_positives < false_positive_limit ~ "heck_yes",
  TRUE ~ "boo"
)) %>%
select(total_trials, alpha_level, false_positive_limit, everything()) %>%
gt_preview()

```

	total_trials	alpha_level	false_positive_limit	false_positives	false_positive_rate	tol_limit_procedure_working?
1	10000	0.1	1000	958	0.0958	heck_yes

Because .0958 is less than .1 (10%) we can say that the tolerance interval procedure is capturing the desired population proportion (95%) with an acceptable false positive rate ($\leq 10\%$). Neat!

Now let's look at the Bayesian analogue.

Bayesian 2-sided Tolerance Interval

Load the Data

The dataset and example used for this section is adapted from [this excellent paper on coverage intervals](#).¹ The values represent measurements of mass fraction of iron made by the National Institute of Standards and Technology (NIST). Units of measurement are percentages, or cg/g.

```

w_tbl <- tibble(w = c(67.43, 66.97, 67.65, 66.84, 67.05, 66.57, 67.16, 68.3,
67.01, 67.07, 67.23, 66.51, 66.46, 67.54, 67.09, 66.77)) %>%
mutate(w_pct = w/100)

```

In preparation for the Bayesian tolerance interval calculations we should run some prior predictive checks. I know Beta(1,1) is a flat prior on mu, but I am not experienced enough with the cauchy to mentally visualize what its parameters mean with respect to the outcome variable. I adjusted the scale parameter for sigma manually and did this visual check each time until I saw a nice cloud of curves of varying but reasonable widths at scale = .025. Bonus - it is so satisfying to make the code for these curves using map2 and dplyr in the prep and then ggplot for the vis - incredibly smooth, readable, and efficient. I remain grateful for these core packages, especially as I start to mess around with other languages.

Prior Predictive Simulations

```

set.seed(2005)

prior_pred_tbl <- tibble(mu = rbeta(300, 1, 1),
  sig = rcauchy(300, 0, 0.025) %>% abs()) %>%
mutate(row_id = row_number()) %>%
select(row_id, everything()) %>%
mutate(plotted_y_data = map2(mu, sig, ~tibble(
  x = seq(0, 1, length.out = 100),
  y = dnorm(x, .x, .y)))) %>%
unnest() %>%
mutate(model = "mu ~ beta(1,1), sigma ~ cauchy(0,0.025)")

prior_pred_tbl %>%
gt_preview()

```

	row_id	mu	sig	x	y	model
1	1	0.1675134	2.18178615	0.00000000	1.823131e-01	mu ~ beta(1,1), sigma ~ cauchy(0,0.025)
2	1	0.1675134	2.18178615	0.01010101	1.823759e-01	mu ~ beta(1,1), sigma ~ cauchy(0,0.025)
3	1	0.1675134	2.18178615	0.02020202	1.824349e-01	mu ~ beta(1,1), sigma ~ cauchy(0,0.025)
4	1	0.1675134	2.18178615	0.03030303	1.824900e-01	mu ~ beta(1,1), sigma ~ cauchy(0,0.025)
5	1	0.1675134	2.18178615	0.04040404	1.825412e-01	mu ~ beta(1,1), sigma ~ cauchy(0,0.025)
6..29999						
30000	300	0.5856645	0.01596839	1.00000000	1.589627e-145	mu ~ beta(1,1), sigma ~ cauchy(0,0.025)

```

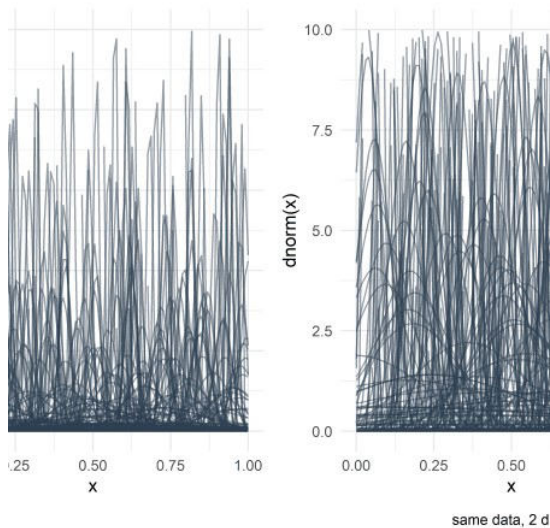
j <- prior_pred_tbl %>%
ggplot(aes(x = x, y = y, group = row_id)) +
geom_line(aes(x, y), alpha = .5, color = "#2c3e50") +
labs(
  x = "x",
  y = "dnorm(x)"
) +
ylim(c(0, 10)) +
theme_minimal()

k <- j +
  ylim(c(0, 50))

(k + j) + plot_annotation(
  title = 'Prior Predictions: n=300 Possible Distributions According to Priors:',
  subtitle = 'mu ~ Beta(1,1), sigma ~ cauchy(0, .025)',
  caption = "same data, 2 different y-axis scales")

```

ns: n=300 Possible Distributions According to Priors:
 $\sigma \sim \text{cauchy}(0, .025)$



Fit the model with brms

Use the priors that were just identified to fit a model in brms and visualize the output. The fuzzy caterpillar is not plotted because its very slow with to render with so many iterations. Don't worry - it's fuzzy.

Non-Informative Priors

```
w_mod <-
  brm(
    data = w_tbl, family = gaussian(),
    w_pct ~ 1,
    prior = c(
      prior(beta(1, 1), class = Intercept),
      prior(cauchy(0, .025), class = sigma)),
    iter = 400000, warmup = 50000, chains = 4, cores = 2,
    seed = 10
  )

w_mod %>% summary()
## Family: gaussian
## Links: mu = identity; sigma = identity
## Formula: w_pct ~ 1
## Data: w_tbl (Number of observations: 16)
## Samples: 4 chains, each with iter = 4e+05; warmup = 50000; thin = 1;
##          total post-warmup samples = 1400000
##
## Population-Level Effects:
##      Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## Intercept      0.67      0.00    0.67    0.67 1.00   858299   757467
##
## Family Specific Parameters:
##      Estimate Est.Error 1-95% CI u-95% CI Rhat Bulk_ESS Tail_ESS
## sigma      0.01      0.00    0.00    0.01 1.00   792388   740061
##
## Samples were drawn using sampling(NUTS). For each parameter, Bulk_ESS
## and Tail_ESS are effective sample size measures, and Rhat is the potential
## scale reduction factor on split chains (at convergence, Rhat = 1).
```

Peek at the posterior tbl for mu and sigma.

```
w_post_tbl <-
  posterior_samples(w_mod) %>%
  select(-lp__) %>%
  rename("mu" = b_Intercept)

w_post_tbl %>%
  gt_preview()

      mu      sigma
1 0.6704184 0.004895014
2 0.6685570 0.005720891
3 0.6708761 0.003932762
4 0.6713238 0.004217306
5 0.6689520 0.005191162
6...1399999
1400000 0.6720105 0.005901202
```

Calculate 2-Sided Bayesian Tolerance Interval from the Posterior - Non-Informative Priors

We'll calculate the 95/90 tolerance limit just like before. The Bayesian analogue involves doing a row-wise calculation for each posterior draw where the equivalent of a tolerance interval is calculated using the standard formula: $x \pm ks$, where x is μ and s is σ . The k factor is taken as the .95 quantile of a unit normal distribution (thus excluding the top and bottom 5% for 90% coverage). Doing this for every row in the posterior produces a distribution of upper and lower tolerance bounds - 1 observation for each credible set of μ and σ . The 5% "confidence" analogue is done by excluding the 2.5% most extreme values from the low end of the lower bound distribution and the 2.5% most extreme values from the high end of the upper bound distribution.

```
bt_lower <- quantile(w_post_tbl$mu - qnorm(0.95) * w_post_tbl$sigma, 0.025)
bt_upper <- quantile(w_post_tbl$mu + qnorm(0.95) * w_post_tbl$sigma, 0.975)
```

```
bt_tbl <- tibble(lower_coverage_limit = bt_lower * 100 %>% round(4),
  upper_coverage_limit = bt_upper * 100 %>% round(4),
  method = as_factor("Bayesian"))
```

```
bt_tbl %>%
  gt_preview()
```

	lower_coverage_limit	upper_coverage_limit	method
1	65.75688	68.44929	Bayesian

So the Bayesian version of a tolerance interval for those relatively vague priors is [65.84, 68.37] after back transforming the percentages to cg/g.

The frequentist version for the same data is calculated as follows with the tolerance package (assuming normality of the experimental data):

Frequentist 2-sided Tolerance Interval

```
fr_tbl <- tidy(normtol.int(w_tbl$w, alpha = .05, P = .9, side = 2, method = "EXACT")) %>%
  select(column, mean) %>%
  pivot_wider(names_from = "column",
    values_from = "mean") %>%
  select('2-sided.lower', '2-sided.upper') %>%
  mutate(method = as_factor("Frequentist")) %>%
  rename(lower_coverage_limit = '2-sided.lower',
    upper_coverage_limit = '2-sided.upper') %>%
  select(method, everything())
```

Combining and comparing to the Bayesian version:

```
fr_tbl %>%
  bind_rows(bt_tbl) %>%
  mutate_if(is.numeric, round, digits = 2) %>%
  mutate(priors = c("non_informative (implied)", "non-informative")) %>%
  gt_preview()
```

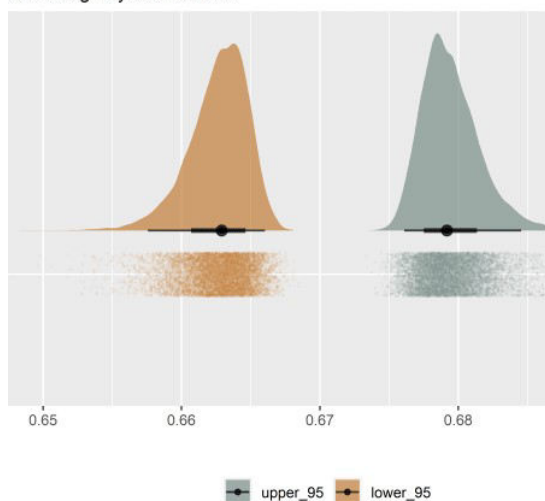
	method	lower_coverage_limit	upper_coverage_limit	priors
1	Frequentist	65.95	68.25	non_informative (implied)
2	Bayesian	65.76	68.45	non-informative

These are quite close! The Bayesian version does remain bit wider. Perhaps with more iterations of MCMC or with slightly less vague priors, the Bayesian version would creep tighter until they converge.

Plotting the upper and lower tolerance limits from the Bayesian version we see that the distribution (upper and lower bounds) are not normal.

```
w_post_tbl %>%
  mutate(upper_95 = qnorm(p = .95, mean = mu, sd = sigma),
    lower_95 = qnorm(p = .05, mean = mu, sd = sigma),
    id = as_factor("posterior")) %>%
  sample_n(size = 10000) %>%
  pivot_longer(cols = c(upper_95, lower_95), names_to = "quantile", values_to = "value") %>%
  mutate(quantile = as_factor(quantile)) %>%
  ggplot(aes(x = value, y = id)) +
  geom_jitter(aes(color = quantile), width = .001, height = .1, size = .3, alpha = .05) +
  stat_halfeye(aes(fill = quantile), alpha = .7, position = position_nudge(y = .2)) +
  scale_color_manual(values = wes_palette("Moonrise2")) +
  scale_fill_manual(values = wes_palette("Moonrise2")) +
  # scale_color_manual(values = c("purple", "limegreen")) +
  # scale_fill_manual(values = c("purple", "limegreen")) +
  labs(
    x = "",
    y = "",
    title = "Distribution of Calculated Upper and Lower Tolerance Limits",
    subtitle = "Calculated using Bayesian Method",
    caption = "10,000 Predictions Sampled from Posterior (400,000 total)"
  ) +
  theme(legend.title = element_blank(),
    legend.position = "bottom")
```

Plot of Calculated Upper and Lower Tolerance Limits Based on Bayesian Method



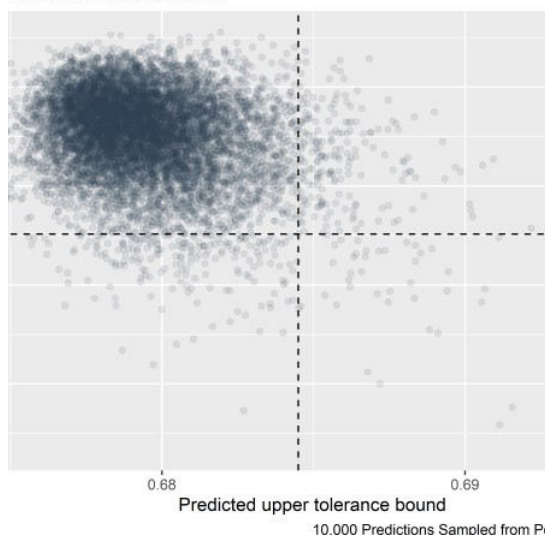
10,000 Predictions Sampled from Posterior

But the upper and lower bounds come in pairs. Plotting the pairs

```
set.seed(2345)
```

```
w_post_tbl %>%
  mutate(upper_95 = qnorm(p = .95, mean = mu, sd = sigma),
         lower_95 = qnorm(p = .05, mean = mu, sd = sigma)) %>%
  sample_n(10000) %>%
  ggplot(aes(y = lower_95, x = upper_95)) +
  geom_point(alpha = .1, color = "#2c3e50") +
  geom_hline(yintercept = bt_lower, linetype = 2) +
  geom_vline(xintercept = bt_upper, linetype = 2) +
  theme_minimal() %>%
  labs(
    x = "Predicted upper tolerance bound",
    y = "Predicted lower tolerance bound",
    title = "Credible Values for Upper and Lower 95% Bounds",
    subtitle = "Based on Credible Posterior Draws",
    caption = "10,000 Predictions Sampled from Posterior (400,000 total)"
  )
```

Values for Upper and Lower 95% Bounds Based on Credible Posterior Draws



10,000 Predictions Sampled from Posterior

Now we can calculate the number of tolerance bounds that are both:

- Above (within) the lower 95% tolerance bound for 90% coverage
- Below (within) the upper 95% tolerance bound for 90% coverage

This number should be less than 500 for a 5% or less false positive rate, since this is a subset of 10,000 draws from the full posterior (which wouldn't display well).

```
set.seed(2345)
```

```
w_post_tbl %>%
  mutate(upper_95 = qnorm(p = .95, mean = mu, sd = sigma),
         lower_95 = qnorm(p = .05, mean = mu, sd = sigma)) %>%
  sample_n(10000) %>%
  filter(lower_95 > bt_lower,
         upper_95 < bt_upper) %>%
  gt_preview()
```

	mu	sigma	upper_95	lower_95
1	0.6734567	0.005124883	0.6818864	0.6650271
2	0.6705619	0.004108660	0.6773200	0.6638037
3	0.6704661	0.005681679	0.6798116	0.6611206
4	0.6709579	0.003404294	0.6765575	0.6653584
5	0.6706663	0.004703547	0.6784029	0.6629296
6...	9541			
9542	0.6713993	0.004295096	0.6784641	0.6643345

9,542 rows means 458 false positives - this is acceptable coverage for the Bayesian interval.